

# Stock Price Prediction: Comparison between ARIMA Model and LSTM

Xizhu Lin

```
In [1]: > import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: > import yfinance as yf
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.stattools import adfuller
import statsmodels.api as sm
from sklearn.preprocessing import MinMaxScaler
```

```
In [3]: > from statsmodels.tsa.arima.model import ARIMA
from sklearn.metrics import mean_squared_error
import pmdarima as pm
```

```
In [4]: > from keras.models import Sequential
from keras.layers import Dense, LSTM
from keras.optimizers import Adam
```

## The Boeing Company (BA) & Zoom Video Communications, Inc. (ZM)

### Import Data

```
In [5]: > start_date = '2020-01-01'
end_date = '2023-05-12'

ba = yf.download("BA", start_date, end_date)['Close']
zm = yf.download("ZM", start_date, end_date)['Close']

[*****100%*****] 1 of 1 completed
[*****100%*****] 1 of 1 completed
```

```
In [6]: > ba = pd.DataFrame(ba)
# check if there is null values in BA Close price
ba[ba.isna().any(axis=1)]
```

Out[6]:

Close
Date

```
In [7]: ▶ # First-order differencing
ba_diff = ba.diff().dropna()
# Log price of BA
ba_lgp = np.log(ba)
# Log Returns of BA
ba_rtn = np.log(ba).diff().dropna()
```

```
In [8]: ▶ zm = pd.DataFrame(zm)
# check if there is null values in ZM Close price
zm[zm.isna().any(axis=1)]
```

Out[8]:

	Close
Date	


```
In [9]: ▶ # First-order differencing
zm_diff = zm.diff().dropna()
# Log price of ZM
zm_lgp = np.log(zm)
# Log Returns of ZM
zm_rtn = np.log(zm).diff().dropna()
```

## EDA

```
In [10]: plt.figure(figsize=(12,8))

plt.subplot(2, 1, 1)
ba['Close'].plot()
plt.grid()
plt.ylabel("Price")
plt.title("The Boeing Company (BA) Share Close Price", fontsize = 16)
plt.subplot(2, 1, 2)
ba_diff['Close'].plot()
plt.grid()
plt.tight_layout(pad=1.5)
plt.ylabel("Daily Price Difference")


plt.savefig("images/Boeing_Company_(BA)_stationary_1")
```



```
In [11]: ▶ plt.figure(figsize=(12,8))

plt.subplot(2, 1, 1)
ba_lgp['Close'].plot()
plt.grid()
plt.ylabel("Price")
plt.title("The Boeing Company (BA) Share Close Log Price", fontsize = 16)
plt.subplot(2, 1, 2)
ba_rtn['Close'].plot()
plt.grid()
plt.tight_layout(pad=1.5)
plt.ylabel("Daily Log Returns")


plt.savefig("images/Boeing_Company_(BA)_stationary_2")
```



```
In [12]: ▶ plt.figure(figsize=(12,8))

plt.subplot(2, 1, 1)
zm['Close'].plot()
plt.grid()
plt.ylabel("Price")
plt.title("Zoom Video Communications, Inc. (ZM) Share Close Price", fontsize = 16)
plt.subplot(2, 1, 2)
zm_rtn['Close'].plot()
plt.grid()
plt.tight_layout(pad=1.5)
plt.ylabel("Daily Log Returns")


plt.savefig("images/Zoom_(ZM)_stationary_1")
```



```
In [13]: ▶ plt.figure(figsize=(12,8))

plt.subplot(2, 1, 1)
zm_lgp['Close'].plot()
plt.grid()
plt.ylabel("Price")
plt.title("Zoom Video Communications, Inc. (ZM) Share Close Log Price", fontsize = 16)
plt.subplot(2, 1, 2)
zm_rtn['Close'].plot()
plt.grid()
plt.tight_layout(pad=1.5)
plt.ylabel("Daily Log Returns")

plt.savefig("images/Zoom_(ZM)_stationary_2")
```




```
In [14]: ▶ df = pd.merge(ba, zm, left_index=True, right_index=True)
df = df.rename(columns={'Close_x': 'Close_BA', 'Close_y': 'Close_ZM'})
```

```
In [15]: ▶ plt.figure(figsize=(12,8))

df['Close_BA'].plot()
df['Close_ZM'].plot()
plt.grid()
plt.ylabel("Price")
plt.title("Boeing (BA) Share Close Price vs Zoom (ZM) Share Close Price", fontsize = 16)
plt.legend(bbox_to_anchor = (1,1))
plt.tight_layout(pad=1.5)

plt.savefig("images/Boeing_and_Zoom")
```



## Price Prediction Model: ARIMA Model and LSTM

### ARIMA Model

#### Boeing (BA)

```
In [16]: ▶ # Get the number of rows to train the model on
training_ba_len = int(np.ceil(len(ba) * .8 ))
# Create the scaled training data set
train_ba = ba[0:int(training_ba_len)]
# First-order differencing
train_ba_diff = train_ba.diff().dropna()
# Log price
train_ba_lgp = np.log(train_ba)
# Log Returns
train_ba_rtn = np.log(train_ba).diff().dropna()
```

## 1. Find d in (p, d, q) for ARIMA model and (ADF) test

```
In [17]: # BA Pirce d=0
result = adfuller(train_ba)
print('ADF Statistic: %f' % result[0])
print('p-value: %f' % result[1])
for key, value in result[4].items():
    print('\t%s: %.3f' % (key, value))
```

ADF Statistic: -3.100203  
p-value: 0.026538  
1%: -3.440  
5%: -2.866  
10%: -2.569

```
In [18]: # BA Pirce d=1
result = adfuller(train_ba_diff)
print('ADF Statistic: %f' % result[0])
print('p-value: %f' % result[1])
for key, value in result[4].items():
    print('\t%s: %.3f' % (key, value))
```

ADF Statistic: -15.435893  
p-value: 0.000000  
1%: -3.440  
5%: -2.866  
10%: -2.569

```
In [19]: # BA Log Pirce d=0
result = adfuller(train_ba_lgp)
print('ADF Statistic: %f' % result[0])
print('p-value: %f' % result[1])
for key, value in result[4].items():
    print('\t%s: %.3f' % (key, value))
```

ADF Statistic: -3.045616  
p-value: 0.030848  
1%: -3.440  
5%: -2.866  
10%: -2.569

```
In [20]: # BA Log Pirce d=1
result = adfuller(train_ba_rtn)
print('ADF Statistic: %f' % result[0])
print('p-value: %f' % result[1])
for key, value in result[4].items():
    print('\t%s: %.3f' % (key, value))
```

ADF Statistic: -8.264118  
p-value: 0.000000  
1%: -3.440  
5%: -2.866  
10%: -2.569



Choose BA Log Pirce d=1, use train\_ba\_lgp (log price of BA) d=1 in ARIMA Model

## 2. Find p and q in (p, d, q) for ARIMA model

### Method 1, Step 1: ACF plot and PACF plot

```
In [21]: ▶ def plot_diagnosticsTight2(data, rtn, title1 = 'APPLE'):

    fig, axes = plt.subplots(1,2,figsize=(15,5))
    axes[0].set_ylabel('ACF', fontsize=15)
    axes[0].set_xlabel('Lag', fontsize=15)
    fig = plot_acf(data,lags=25,zero=True, ax = axes[0], title= 'ACF ' + title1, use_vlines = True)
    axes[1].set_ylabel('ACF', fontsize=15)
    axes[1].set_xlabel('Lag', fontsize=15)
    fig = plot_acf(rtn,lags=25,zero=True, ax = axes[1], title= 'ACF Log Returns ' + title1, use_vli
    plt.tight_layout()
    plt.savefig("ACF2.png", dpi = 300)
```

```
In [22]: ▶ def plot_ACF_PACF(data, lag_num, company_name):
    fig, axes = plt.subplots(1,2,figsize=(15,5))


    fig = plot_acf(data.values, lags = lag_num, zero=False, ax = axes[0],
                    title= 'ACF of ' + company_name + ' log prices')

    fig = plot_pacf(data.values, lags = lag_num, zero=False, ax = axes[1],
                    title= 'PACF of ' + company_name + ' log prices')

    # plt.show()
    plt.savefig("images/ACF PACF" + company_name, dpi = 300)
```

```
In [23]: ▶ plot_ACF_PACF(train_ba_lgp, 25, 'Boeing (BA)')

#plt.savefig("images/Boeing_Company_(BA)_ACF_PACF")
```



**p=2, q=0**

```
In [24]: p, d, q = 2, 1, 0
ar_model_ba = ARIMA(train_ba_lgp['Close'], order=(p, d, q))
ar_model_ba = ar_model_ba.fit()
ar_model_ba.summary()
```

C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.  
 self.\_init\_dates(dates, freq)  
C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.  
 self.\_init\_dates(dates, freq)  
C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.  
 self.\_init\_dates(dates, freq)

Out[24]: SARIMAX Results

<b>Dep. Variable:</b>	Close	<b>No. Observations:</b>	677
<b>Model:</b>	ARIMA(2, 1, 0)	<b>Log Likelihood</b>	1236.277
<b>Date:</b>	Fri, 15 Dec 2023	<b>AIC</b>	-2466.554
<b>Time:</b>	20:01:30	<b>BIC</b>	-2453.005
<b>Sample:</b>	0	<b>HQIC</b>	-2461.308
			- 677
<b>Covariance Type:</b>	opg		

	coef	std err	z	P> z	[0.025	0.975]
<b>ar.L1</b>	0.0636	0.018	3.577	0.000	0.029	0.098
<b>ar.L2</b>	0.1626	0.017	9.523	0.000	0.129	0.196
<b>sigma2</b>	0.0015	4.04e-05	37.406	0.000	0.001	0.002

<b>Ljung-Box (L1) (Q):</b>	0.05	<b>Jarque-Bera (JB):</b>	1515.64
<b>Prob(Q):</b>	0.82	<b>Prob(JB):</b>	0.00
<b>Heteroskedasticity (H):</b>	0.26	<b>Skew:</b>	-0.08
<b>Prob(H) (two-sided):</b>	0.00	<b>Kurtosis:</b>	10.33

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

## 2. Find p and q in (p, d, q) for ARIMA model

### Method 1, Step 2: p-value test

```
In [25]: ▶ # Set the significance level (alpha)
alpha = 0.05

p_values = ar_model_ba.pvalues

significant_predictors = p_values[p_values <= alpha]

print("Significant Predictors:")
print(significant_predictors)
```

```
Significant Predictors:
ar.L1      3.471556e-04
ar.L2      1.684478e-21
sigma2     3.083713e-306
dtype: float64
```

## 2. Find p and q in (p, d, q) for ARIMA model

### Method 2: grid search for min AIC and min BIC

```
In [26]: ▶ def searchARIMA(data, d, max_p, max_q):
    results = []

    aic = np.zeros((max_p, max_q))
    bic = np.zeros((max_p, max_q))

    for i in range(max_p):
        for j in range(max_q):
            model = ARIMA(data, order=(i, d, j))
            model = ARIMA(data, order=(i, d, j))
            res = model.fit()
            aic[i, j] = res.aic
            bic[i, j] = res.bic

            result_dict = {
                'p': i,
                'q': j,
                'aic': aic[i, j],
                'bic': bic[i, j]
            }
            results.append(result_dict)
            # print('p:', i, ' q:', j, ' aic:', aic[i, j], ' bic:', bic[i, j])

    result_df = pd.DataFrame(results)
    return result_df
```

```
In [27]: result_df = searchARIMA(train_ba_lgp.values, d = 1, max_p = 10, max_q = 10)
result_df
```

	p	q	aic	bic
0	0	0	-2448.614890	-2444.098697
1	0	1	-2449.534203	-2440.501817
2	0	2	-2469.495778	-2455.947199
3	0	3	-2467.608602	-2449.543830
4	0	4	-2465.744453	-2443.163488
...	...	...	...	...
95	9	5	-2460.106052	-2392.363156
96	9	6	-2459.101628	-2386.842539
97	9	7	-2459.114381	-2382.339099
98	9	8	-2457.036662	-2375.745187
99	9	9	-2456.636837	-2370.829169

100 rows x 4 columns

```
In [28]: min_aic_bic_row = result_df[(result_df['aic'] == result_df['aic'].min() |
                                     (result_df['bic'] == result_df['bic'].min()))]
min_aic_bic_row
```

Out[28]:

	p	q	aic	bic
2	0	2	-2469.495778	-2455.947199
46	4	6	-2470.604216	-2420.926092

```
In [29]: pm.arma.auto_arma(train_ba_lgp)
```

Out[29]:

	ARIMA
ARIMA(0,1,2)(0,0,0)[0]	

## In general, choose $p=2$ , $q=0$

```
In [30]: >>> p, d, q = 2, 1, 0
ar_model_ba = ARIMA(train_ba_lgp['Close'], order=(p, d, q))
ar_model_ba = ar_model_ba.fit()
ar_model_ba.summary()
```

```
C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.
  self._init_dates(dates, freq)
C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.
  self._init_dates(dates, freq)
C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.
  self._init_dates(dates, freq)
```

Out[30]: SARIMAX Results

<b>Dep. Variable:</b>	Close	<b>No. Observations:</b>	677
<b>Model:</b>	ARIMA(2, 1, 0)	<b>Log Likelihood</b>	1236.277
<b>Date:</b>	Fri, 15 Dec 2023	<b>AIC</b>	-2466.554
<b>Time:</b>	20:04:53	<b>BIC</b>	-2453.005
<b>Sample:</b>	0	<b>HQIC</b>	-2461.308
	- 677		
<b>Covariance Type:</b>	opg		

	coef	std err	z	P> z	[0.025	0.975]
<b>ar.L1</b>	0.0636	0.018	3.577	0.000	0.029	0.098
<b>ar.L2</b>	0.1626	0.017	9.523	0.000	0.129	0.196
<b>sigma2</b>	0.0015	4.04e-05	37.406	0.000	0.001	0.002

<b>Ljung-Box (L1) (Q):</b>	0.05	<b>Jarque-Bera (JB):</b>	1515.64
<b>Prob(Q):</b>	0.82	<b>Prob(JB):</b>	0.00
<b>Heteroskedasticity (H):</b>	0.26	<b>Skew:</b>	-0.08
<b>Prob(H) (two-sided):</b>	0.00	<b>Kurtosis:</b>	10.33

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [31]: >>> # Set the significance level (alpha)
alpha = 0.05

p_values = ar_model_ba.pvalues

significant_predictors = p_values[p_values <= alpha]

print("Significant Predictors:")
print(significant_predictors)
```

```
Significant Predictors:
ar.L1      3.471556e-04
ar.L2      1.684478e-21
sigma2     3.083713e-306
dtype: float64
```

```
In [32]: predicted_original_train_ba = ar_model_ba.predict(start=train_ba_lgp.index[0],
                                                         end=train_ba_lgp.index[-1],
                                                         dynamic=False)
predicted_original_train_ba = np.exp(predicted_original_train_ba)
#predicted_original_ba = np.array(predicted_original_ba)
```

```
In [33]: predicted_original_train_ba = pd.DataFrame(predicted_original_train_ba)

train_val_ba = ba[:int(training_ba_len)]
train_val_ba = pd.merge(train_val_ba, predicted_original_train_ba,
                        left_index=True, right_index=True)
train_val_ba = train_val_ba.rename(columns={'predicted_mean': 'Prediction'})
train_val_ba
```

Out[33]:

	Close	Prediction
Date		
2020-01-02	333.320007	1.000000
2020-01-03	332.760010	333.320008
2020-01-06	333.739990	332.717537
2020-01-07	337.279999	333.711147
2020-01-08	331.369995	337.667707
...	...	...
2022-09-01	153.660004	159.616873
2022-09-02	151.820007	152.947757
2022-09-06	152.389999	150.671611
2022-09-07	155.949997	152.128066
2022-09-08	157.789993	156.274284

677 rows × 2 columns

```
In [34]: train_mse = mean_squared_error(train_val_ba['Close'], train_val_ba['Prediction'])
train_rmse = np.sqrt(train_mse)
print("The MSE of the Trainning set is", train_mse)
print("The RMSE of the Trainning set is", train_rmse)
```

The MSE of the Trainning set is 210.71958656497665  
The RMSE of the Trainning set is 14.516183608820077

```
In [35]: # Create the scaled training data set
test_ba = ba[int(training_ba_len):]
# Log price
test_ba_lgp = np.log(test_ba)

forecast_ba = ar_model_ba.get_forecast(steps=len(test_ba_lgp))

predicted_values_ba = forecast_ba.predicted_mean
predicted_original_ba = np.exp(predicted_values_ba)
predicted_original_ba = np.array(predicted_original_ba)
```

C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:836: ValueWarning: No supported index is available. Prediction results will be given with an integer index beginning at `start`.

```
    return get_prediction_index(
```

```
In [36]: predicted_original_ba = pd.DataFrame(index=test_ba_lgp.index, columns=['Prediction'],
                                             data=predicted_original_ba)

test_ba = pd.merge(test_ba, predicted_original_ba,
                    left_index=True, right_index=True)
test_ba
```

Out[36]:


	Close	Prediction
Date		
2022-09-09	157.520004	158.501649
2022-09-12	158.720001	158.849633
2022-09-13	147.309998	158.988054
2022-09-14	149.259995	159.053558
2022-09-15	149.779999	159.080248
...	...	...
2023-05-05	198.339996	159.101817
2023-05-08	197.259995	159.101817
2023-05-09	201.880005	159.101817
2023-05-10	200.839996	159.101817
2023-05-11	201.839996	159.101817

169 rows × 2 columns

```
In [37]: test_mse = mean_squared_error(test_ba['Close'], test_ba['Prediction'])
test_rmse = np.sqrt(test_mse)
print("The MSE of the Test set is", test_mse)
print("The RMSE of the Test set is", test_rmse)
```

The MSE of the Test set is 1475.7237845602306  
The RMSE of the Test set is 38.41515045604053

```
In [38]: # Visualize the data
plt.figure(figsize=(16,6))
plt.title('ARIMA Model (2,1,0) - Boeing (BA)')
plt.xlabel('Date', fontsize=18)
plt.ylabel('Close Price USD ($)', fontsize=18)
plt.plot(train_val_ba[['Close', 'Prediction']])
plt.plot(test_ba[['Close', 'Prediction']])
plt.legend(['Train_Close', 'Train_Prediction', 'Test_Close', 'Test_Prediction'], loc='upper right')
plt.show()
plt.savefig("images/Boeing_Company_(BA)_ARIMA_1")
```



```
In [39]: p, d, q = 2, 0, 0
ar_model_ba = ARIMA(train_ba_lgp['Close'], order=(p, d, q))
ar_model_ba = ar_model_ba.fit()
ar_model_ba.summary()
```

C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.  
 self.\_init\_dates(dates, freq)  
C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.  
 self.\_init\_dates(dates, freq)  
C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.  
 self.\_init\_dates(dates, freq)

Out[39]:

SARIMAX Results

<b>Dep. Variable:</b>	Close	<b>No. Observations:</b>	677
<b>Model:</b>	ARIMA(2, 0, 0)	<b>Log Likelihood</b>	1229.167
<b>Date:</b>	Fri, 15 Dec 2023	<b>AIC</b>	-2450.334
<b>Time:</b>	20:04:54	<b>BIC</b>	-2432.263
<b>Sample:</b>	0	<b>HQIC</b>	-2443.338
	- 677		
<b>Covariance Type:</b>	opg		

	coef	std err	z	P> z	[0.025	0.975]
<b>const</b>	5.3002	0.089	59.265	0.000	5.125	5.476
<b>ar.L1</b>	1.0705	0.019	55.603	0.000	1.033	1.108
<b>ar.L2</b>	-0.0832	0.018	-4.549	0.000	-0.119	-0.047
<b>sigma2</b>	0.0015	3.61e-05	42.695	0.000	0.001	0.002

<b>Ljung-Box (L1) (Q):</b>	0.25	<b>Jarque-Bera (JB):</b>	2340.08
<b>Prob(Q):</b>	0.62	<b>Prob(JB):</b>	0.00
<b>Heteroskedasticity (H):</b>	0.26	<b>Skew:</b>	-0.54
<b>Prob(H) (two-sided):</b>	0.00	<b>Kurtosis:</b>	12.04

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [40]: # Set the significance level (alpha)
alpha = 0.05

p_values = ar_model_ba.pvalues

significant_predictors = p_values[p_values <= alpha]

print("Significant Predictors:")
print(significant_predictors)
```

Significant Predictors:  
 const 0.000000  
 ar.L1 0.000000  
 ar.L2 0.000005  
 sigma2 0.000000  
 dtype: float64



```
In [41]: ► predicted_original_train_ba = ar_model_ba.predict(start=train_ba_lgp.index[0],
                                                             end=train_ba_lgp.index[-1],
                                                             dynamic=False)
predicted_original_train_ba = np.exp(predicted_original_train_ba)
#predicted_original_ba = np.array(predicted_original_ba)
predicted_original_train_ba = pd.DataFrame(predicted_original_train_ba)

train_val_ba = ba[:int(training_ba_len)]
train_val_ba = pd.merge(train_val_ba, predicted_original_train_ba,
                        left_index=True, right_index=True)
train_val_ba = train_val_ba.rename(columns={'predicted_mean': 'Prediction'})
train_val_ba
```

Out[41]:

	Close	Prediction
Date		
2020-01-02	333.320007	200.386139
2020-01-03	332.760010	331.342606
2020-01-06	333.739990	330.583051
2020-01-07	337.279999	331.671768
2020-01-08	331.369995	335.357282
...	...	...
2022-09-01	153.660004	160.541874
2022-09-02	151.820007	153.639991
2022-09-06	152.389999	152.202058
2022-09-07	155.949997	152.967068
2022-09-08	157.789993	156.746797

677 rows × 2 columns

```
In [42]: ► train_mse = mean_squared_error(train_val_ba['Close'], train_val_ba['Prediction'])
train_rmse = np.sqrt(train_mse)
print("The MSE of the Trainning set is", train_mse)
print("The RMSE of the Trainning set is", train_rmse)
```

The MSE of the Trainning set is 74.27961633410533  
The RMSE of the Trainning set is 8.618562312480273

```
In [43]: # Create the scaled training data set
test_ba = ba[int(training_ba_len):]
# Log price
test_ba_lgp = np.log(test_ba)

forecast_ba = ar_model_ba.get_forecast(steps=len(test_ba_lgp))

predicted_values_ba = forecast_ba.predicted_mean
predicted_original_ba = np.exp(predicted_values_ba)
predicted_original_ba = np.array(predicted_original_ba)
predicted_original_ba = pd.DataFrame(index=test_ba_lgp.index, columns=['Prediction'],
                                     data=predicted_original_ba)

test_ba = pd.merge(test_ba, predicted_original_ba,
                   left_index=True, right_index=True)
test_ba
```

C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:836: ValueWarning: No supported index is available. Prediction results will be given with an integer index beginning at `start`.

```
return get_prediction_index(
```

Out[43]:

	Close	Prediction
Date		
2022-09-09	157.520004	158.422824
2022-09-12	158.720001	158.947933
2022-09-13	147.309998	159.458901
2022-09-14	149.259995	159.963682
2022-09-15	149.779999	160.462985
...	...	...
2023-05-05	198.339996	195.649142
2023-05-08	197.259995	195.713897
2023-05-09	201.880005	195.777777
2023-05-10	200.839996	195.840793
2023-05-11	201.839996	195.902958


169 rows × 2 columns

```
In [44]: test_mse = mean_squared_error(test_ba['Close'], test_ba['Prediction'])
test_rmse = np.sqrt(test_mse)
print("The MSE of the Test set is", test_mse)
print("The RMSE of the Test set is", test_rmse)
```

The MSE of the Test set is 382.83536670614063

The RMSE of the Test set is 19.56617915450384

```
In [45]: # Visualize the data
plt.figure(figsize=(16,6))
plt.title('ARIMA Model (2,0,0) - Boeing (BA)')
plt.xlabel('Date', fontsize=18)
plt.ylabel('Close Price USD ($)', fontsize=18)
plt.plot(train_val_ba[['Close', 'Prediction']])
plt.plot(test_ba[['Close', 'Prediction']])
plt.legend(['Train_Close', 'Train_Prediction', 'Test_Close', 'Test_Prediction'], loc='upper right')
plt.show()
plt.savefig("images/Boeing_Company_(BA)_ARIMA_2")
```



```
In [46]: pm.arma.auto_arma(train_ba_lgp)
```

Out[46]:

```
ARIMA
ARIMA(0,1,2)(0,0,0)[0]
```

```
In [47]: p, d, q = 0, 1, 2
ar_model_ba = ARIMA(train_ba_lgp['Close'], order=(p, d, q))
ar_model_ba = ar_model_ba.fit()
ar_model_ba.summary()
```

C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.  
 self.\_init\_dates(dates, freq)  
C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.  
 self.\_init\_dates(dates, freq)  
C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.  
 self.\_init\_dates(dates, freq)

Out[47]:

SARIMAX Results

<b>Dep. Variable:</b>	Close	<b>No. Observations:</b>	677
<b>Model:</b>	ARIMA(0, 1, 2)	<b>Log Likelihood</b>	1237.748
<b>Date:</b>	Fri, 15 Dec 2023	<b>AIC</b>	-2469.496
<b>Time:</b>	20:05:01	<b>BIC</b>	-2455.947
<b>Sample:</b>	0	<b>HQIC</b>	-2464.250
	- 677		
<b>Covariance Type:</b>	opg		

	coef	std err	z	P> z	[0.025	0.975]
<b>ma.L1</b>	0.0720	0.017	4.190	0.000	0.038	0.106
<b>ma.L2</b>	0.1834	0.020	9.382	0.000	0.145	0.222
<b>sigma2</b>	0.0015	4.19e-05	35.863	0.000	0.001	0.002

<b>Ljung-Box (L1) (Q):</b>	0.00	<b>Jarque-Bera (JB):</b>	1446.36
<b>Prob(Q):</b>	0.95	<b>Prob(JB):</b>	0.00
<b>Heteroskedasticity (H):</b>	0.26	<b>Skew:</b>	-0.11
<b>Prob(H) (two-sided):</b>	0.00	<b>Kurtosis:</b>	10.16

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [48]: # Set the significance level (alpha)
alpha = 0.05

p_values = ar_model_ba.pvalues

significant_predictors = p_values[p_values <= alpha]

print("Significant Predictors:")
print(significant_predictors)
```

Significant Predictors:  
 ma.L1 2.794185e-05  
 ma.L2 6.489390e-21  
 sigma2 1.167885e-281  
 dtype: float64

```
In [49]: predicted_original_train_ba = ar_model_ba.predict(start=train_ba_lgp.index[0],
                                                         end=train_ba_lgp.index[-1],
                                                         dynamic=False)
predicted_original_train_ba = np.exp(predicted_original_train_ba)
#predicted_original_ba = np.array(predicted_original_ba)
predicted_original_train_ba = pd.DataFrame(predicted_original_train_ba)

train_val_ba = ba[:int(training_ba_len)]
train_val_ba = pd.merge(train_val_ba, predicted_original_train_ba,
                        left_index=True, right_index=True)
train_val_ba = train_val_ba.rename(columns={'predicted_mean': 'Prediction'})
train_val_ba
```

Out[49]:

	Close	Prediction
Date		
2020-01-02	333.320007	1.000000
2020-01-03	332.760010	333.320008
2020-01-06	333.739990	332.714107
2020-01-07	337.279999	333.710779
2020-01-08	331.369995	337.725523
...	...	...
2022-09-01	153.660004	159.728619
2022-09-02	151.820007	152.904828
2022-09-06	152.389999	150.667790
2022-09-07	155.949997	152.315716
2022-09-08	157.789993	156.541106

677 rows × 2 columns

```
In [50]: train_mse = mean_squared_error(train_val_ba['Close'], train_val_ba['Prediction'])
train_rmse = np.sqrt(train_mse)
print("The MSE of the Trainning set is", train_mse)
print("The RMSE of the Trainning set is", train_rmse)
```

The MSE of the Trainning set is 210.62604708648428  
The RMSE of the Trainning set is 14.512961347929108

```
In [51]: # Create the scaled training data set
test_ba = ba[int(training_ba_len):]
# Log price
test_ba_lgp = np.log(test_ba)

forecast_ba = ar_model_ba.get_forecast(steps=len(test_ba_lgp))

predicted_values_ba = forecast_ba.predicted_mean
predicted_original_ba = np.exp(predicted_values_ba)
predicted_original_ba = np.array(predicted_original_ba)
predicted_original_ba = pd.DataFrame(index=test_ba_lgp.index, columns=['Prediction'],
                                     data=predicted_original_ba)

test_ba = pd.merge(test_ba, predicted_original_ba,
                   left_index=True, right_index=True)
test_ba
```

C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:836: ValueWarning: No supported index is available. Prediction results will be given with an integer index beginning at `start`.

```
return get_prediction_index(
```

Out[51]:


	Close	Prediction
Date		
2022-09-09	157.520004	158.564732
2022-09-12	158.720001	158.796043
2022-09-13	147.309998	158.796043
2022-09-14	149.259995	158.796043
2022-09-15	149.779999	158.796043
...	...	...
2023-05-05	198.339996	158.796043
2023-05-08	197.259995	158.796043
2023-05-09	201.880005	158.796043
2023-05-10	200.839996	158.796043
2023-05-11	201.839996	158.796043

169 rows × 2 columns

```
In [52]: test_mse = mean_squared_error(test_ba['Close'], test_ba['Prediction'])
test_rmse = np.sqrt(test_mse)
print("The MSE of the Test set is", test_mse)
print("The RMSE of the Test set is", test_rmse)
```

The MSE of the Test set is 1491.442976445892  
The RMSE of the Test set is 38.61920476195609

```
In [53]: # Visualize the data
plt.figure(figsize=(16,6))
plt.title('ARIMA Model (0,1,2) - Boeing (BA)')
plt.xlabel('Date', fontsize=18)
plt.ylabel('Close Price USD ($)', fontsize=18)
plt.plot(train_val_ba[['Close', 'Prediction']])
plt.plot(test_ba[['Close', 'Prediction']])
plt.legend(['Train_Close', 'Train_Prediction', 'Test_Close', 'Test_Prediction'], loc='upper right')
#plt.show()
plt.savefig("images/Boeing_Company_(BA)_ARIMA_3")
```



```
In [54]: p, d, q = 4, 1, 6
ar_model_ba = ARIMA(train_ba_lgp['Close'], order=(p, d, q))
ar_model_ba = ar_model_ba.fit()
ar_model_ba.summary()
```

```
C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.
  self._init_dates(dates, freq)
C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.
  self._init_dates(dates, freq)
C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.
  self._init_dates(dates, freq)
C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\base\model.py:607: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
  warnings.warn("Maximum Likelihood optimization failed to "
```

Out[54]: SARIMAX Results

<b>Dep. Variable:</b>	Close	<b>No. Observations:</b>	677
<b>Model:</b>	ARIMA(4, 1, 6)	<b>Log Likelihood</b>	1246.302
<b>Date:</b>	Fri, 15 Dec 2023	<b>AIC</b>	-2470.604
<b>Time:</b>	20:05:03	<b>BIC</b>	-2420.926
<b>Sample:</b>	0	<b>HQIC</b>	-2451.370
	- 677		
<b>Covariance Type:</b>	opg		

	coef	std err	z	P> z	[0.025	0.975]
<b>ar.L1</b>	-0.5023	0.071	-7.026	0.000	-0.642	-0.362
<b>ar.L2</b>	0.6586	0.077	8.539	0.000	0.507	0.810
<b>ar.L3</b>	-0.4637	0.066	-7.073	0.000	-0.592	-0.335
<b>ar.L4</b>	-0.8479	0.062	-13.669	0.000	-0.969	-0.726
<b>ma.L1</b>	0.5777	0.074	7.776	0.000	0.432	0.723
<b>ma.L2</b>	-0.4433	0.082	-5.402	0.000	-0.604	-0.282
<b>ma.L3</b>	0.5144	0.069	7.415	0.000	0.378	0.650
<b>ma.L4</b>	0.7332	0.073	10.104	0.000	0.591	0.875
<b>ma.L5</b>	0.1076	0.035	3.058	0.002	0.039	0.177
<b>ma.L6</b>	0.1051	0.037	2.819	0.005	0.032	0.178
<b>sigma2</b>	0.0015	4.67e-05	31.334	0.000	0.001	0.002

<b>Ljung-Box (L1) (Q):</b>	0.00	<b>Jarque-Bera (JB):</b>	1253.46
<b>Prob(Q):</b>	0.99	<b>Prob(JB):</b>	0.00
<b>Heteroskedasticity (H):</b>	0.27	<b>Skew:</b>	-0.15
<b>Prob(H) (two-sided):</b>	0.00	<b>Kurtosis:</b>	9.66

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).



```
In [55]: # Set the significance level (alpha)
alpha = 0.05

p_values = ar_model_ba.pvalues

significant_predictors = p_values[p_values <= alpha]

print("Significant Predictors:")
print(significant_predictors)
```

```
Significant Predictors:
ar.L1      2.127745e-12
ar.L2      1.358088e-17
ar.L3      1.515523e-12
ar.L4      1.563808e-42
ma.L1      7.455934e-15
ma.L2      6.600023e-08
ma.L3      1.212720e-13
ma.L4      5.302499e-24
ma.L5      2.226047e-03
ma.L6      4.812607e-03
sigma2     1.630318e-215
dtype: float64
```

```
In [56]: predicted_original_train_ba = ar_model_ba.predict(start=train_ba_lgp.index[0],
                                                            end=train_ba_lgp.index[-1],
                                                            dynamic=False)
predicted_original_train_ba = np.exp(predicted_original_train_ba)
#predicted_original_ba = np.array(predicted_original_ba)
predicted_original_train_ba = pd.DataFrame(predicted_original_train_ba)

train_val_ba = ba[:int(training_ba_len)]
train_val_ba = pd.merge(train_val_ba, predicted_original_train_ba,
                        left_index=True, right_index=True)
train_val_ba = train_val_ba.rename(columns={'predicted_mean': 'Prediction'})
train_val_ba
```

Out[56]:

	Close	Prediction
Date		
2020-01-02	333.320007	1.000000
2020-01-03	332.760010	333.320008
2020-01-06	333.739990	332.718327
2020-01-07	337.279999	333.703758
2020-01-08	331.369995	337.720118
...	...	...
2022-09-01	153.660004	159.292708
2022-09-02	151.820007	153.214733
2022-09-06	152.389999	151.324491
2022-09-07	155.949997	152.463883
2022-09-08	157.789993	157.347112

677 rows × 2 columns

```
In [57]: train_mse = mean_squared_error(train_val_ba['Close'], train_val_ba['Prediction'])
train_rmse = np.sqrt(train_mse)
print("The MSE of the Training set is", train_mse)
print("The RMSE of the Training set is", train_rmse)
```

```
The MSE of the Training set is 209.78057389758536
The RMSE of the Training set is 14.483803847663271
```

```
In [58]: # Create the scaled training data set
test_ba = ba[int(training_ba_len):]
# Log price
test_ba_lgp = np.log(test_ba)

forecast_ba = ar_model_ba.get_forecast(steps=len(test_ba_lgp))

predicted_values_ba = forecast_ba.predicted_mean
predicted_original_ba = np.exp(predicted_values_ba)
predicted_original_ba = np.array(predicted_original_ba)
predicted_original_ba = pd.DataFrame(index=test_ba_lgp.index, columns=['Prediction'],
                                     data=predicted_original_ba)

test_ba = pd.merge(test_ba, predicted_original_ba,
                   left_index=True, right_index=True)
test_ba
```

C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:836: ValueWarning: No supported index is available. Prediction results will be given with an integer index beginning at `start`.


```
    return get_prediction_index(
```

Out[58]:

	Close	Prediction
Date		
2022-09-09	157.520004	157.986055
2022-09-12	158.720001	158.615432
2022-09-13	147.309998	157.291470
2022-09-14	149.259995	157.533607
2022-09-15	149.779999	156.511250
...	...	...
2023-05-05	198.339996	157.562173
2023-05-08	197.259995	157.563195
2023-05-09	201.880005	157.563325
2023-05-10	200.839996	157.564286
2023-05-11	201.839996	157.563254

169 rows × 2 columns

```
In [59]: # Visualize the data
plt.figure(figsize=(16,6))
plt.title('ARIMA Model (4,1,6) - Boeing (BA)')
plt.xlabel('Date', fontsize=18)
plt.ylabel('Close Price USD ($)', fontsize=18)
plt.plot(train_val_ba[['Close', 'Prediction']])
plt.plot(test_ba[['Close', 'Prediction']])
plt.legend(['Train_Close', 'Train_Prediction', 'Test_Close', 'Test_Prediction'], loc='upper right')
plt.show()
plt.savefig("images/Boeing_Company_(BA)_ARIMA_4")
```



## Zoom (ZM)

```
In [60]: # Get the number of rows to train the model on
training_zm_len = int(np.ceil(len(zm) * .8))
# Create the scaled training data set
train_zm = zm[0:int(training_zm_len)]
# First-order differencing
train_zm_diff = train_zm.diff().dropna()
# Log price
train_zm_lgp = np.log(train_zm)
# Log Returns
train_zm_rtn = np.log(train_zm).diff().dropna()
```

## 1. Find d in (p, d, q) for ARIMA model and (ADF) test

```
In [61]: # ZM Pirce d=0
result = adfuller(train_zm)
print('ADF Statistic: %f' % result[0])
print('p-value: %f' % result[1])
for key, value in result[4].items():
    print('\t%s: %.3f' % (key, value))
```

```
ADF Statistic: -1.091567
p-value: 0.718423
1%: -3.440
5%: -2.866
10%: -2.569
```

```
In [62]: # ZM Pirce d=1
result = adfuller(train_zm_diff)
print('ADF Statistic: %f' % result[0])
print('p-value: %f' % result[1])
for key, value in result[4].items():
    print('\t%s: %.3f' % (key, value))
```

```
ADF Statistic: -17.269873
p-value: 0.000000
1%: -3.440
5%: -2.866
10%: -2.569
```

```
In [63]: # ZM Log Pirce d=0
result = adfuller(train_zm_lgp)
print('ADF Statistic: %f' % result[0])
print('p-value: %f' % result[1])
for key, value in result[4].items():
    print('\t%s: %.3f' % (key, value))
```

```
ADF Statistic: -1.248085
p-value: 0.652576
1%: -3.440
5%: -2.866
10%: -2.569
```

```
In [64]: # ZM Log Pirce d=1
result = adfuller(train_zm_rtn)
print('ADF Statistic: %f' % result[0])
print('p-value: %f' % result[1])
for key, value in result[4].items():
    print('\t%s: %.3f' % (key, value))
```


```
ADF Statistic: -26.131050
p-value: 0.000000
1%: -3.440
5%: -2.866
10%: -2.569
```

Choose ZM Log Pirce d=1, use train\_zm\_lgp (log price of ZM) d=1 in ARIMA Model

2. Find p and q in (p, d, q) for ARIMA model

Method 1, Step 1: ACF plot and PACF plot

```
In [65]: plot_ACF_PACF(train_zm_lgp, 25, 'Zoom (ZM)')  
#plt.savefig("images/Zoom_(ZM)_ACF_PACF")
```



**p=0, q=0**

```
In [66]: > p, d, q = 0, 1, 0
ar_model_zm = ARIMA(train_zm_lgp['Close'], order=(p, d, q))
ar_model_zm = ar_model_zm.fit()
ar_model_zm.summary()
```

C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.  
 self.\_init\_dates(dates, freq)  
C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.  
 self.\_init\_dates(dates, freq)  
C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.  
 self.\_init\_dates(dates, freq)

Out[66]: SARIMAX Results

<b>Dep. Variable:</b>	Close	<b>No. Observations:</b>	677
<b>Model:</b>	ARIMA(0, 1, 0)	<b>Log Likelihood</b>	1160.345
<b>Date:</b>	Fri, 15 Dec 2023	<b>AIC</b>	-2318.689
<b>Time:</b>	20:05:06	<b>BIC</b>	-2314.173
<b>Sample:</b>	0	<b>HQIC</b>	-2316.941
	- 677		
<b>Covariance Type:</b>	opg		

	coef	std err	z	P> z	[0.025	0.975]
<b>sigma2</b>	0.0019	4.69e-05	40.301	0.000	0.002	0.002

<b>Ljung-Box (L1) (Q):</b>	0.03	<b>Jarque-Bera (JB):</b>	1646.47
<b>Prob(Q):</b>	0.85	<b>Prob(JB):</b>	0.00
<b>Heteroskedasticity (H):</b>	0.59	<b>Skew:</b>	0.43
<b>Prob(H) (two-sided):</b>	0.00	<b>Kurtosis:</b>	10.60

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

## Method 1, Step 2: p-value test

```
In [67]: > # Set the significance level (alpha)
alpha = 0.05

p_values = ar_model_zm.pvalues

significant_predictors = p_values[p_values <= alpha]

print("Significant Predictors:")
print(significant_predictors)
```

Significant Predictors:  
 sigma2 0.0  
 dtype: float64

## 2. Find p and q in (p, d, q) for ARIMA model

### Method 2: grid search for min AIC and min BIC

```
In [68]: result_df = searchARIMA(train_zm_lgp.values, d = 1, max_p = 10, max_q = 10)
result_df
```

	p	q	aic	bic
0	0	0	-2318.689199	-2314.173006
1	0	1	-2316.725490	-2307.693104
2	0	2	-2315.648769	-2302.100190
3	0	3	-2313.926836	-2295.862063
4	0	4	-2312.859576	-2290.278611
...	...	...	...	...
95	9	5	-2296.018761	-2228.275865
96	9	6	-2298.606085	-2226.346996
97	9	7	-2293.419767	-2216.644485
98	9	8	-2291.103879	-2209.812404
99	9	9	-2294.160153	-2208.352485

100 rows × 4 columns

```
In [69]: min_aic_bic_row = result_df[(result_df['aic'] == result_df['aic'].min()) |
min_aic_bic_row
                                             (result_df['bic'] == result_df['bic'].min())]
```

Out[69]:

	p	q	aic	bic
0	0	0	-2318.689199	-2314.173006

```
In [70]: pm.arma.auto_arma(train_zm_lgp)
```

Out[70]:

```
ARIMA
ARIMA(4,2,1)(0,0,0)[0]
```

**In general, choose  $p=0$ ,  $q=0$** 

```
In [71]: >>> p, d, q = 0, 1, 0
ar_model_zm = ARIMA(train_zm_lgp['Close'], order=(p, d, q))
ar_model_zm = ar_model_zm.fit()
ar_model_zm.summary()
```

```
C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.
  self._init_dates(dates, freq)
C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.
  self._init_dates(dates, freq)
C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.
  self._init_dates(dates, freq)
```

Out[71]: SARIMAX Results

<b>Dep. Variable:</b>	Close	<b>No. Observations:</b>	677
<b>Model:</b>	ARIMA(0, 1, 0)	<b>Log Likelihood</b>	1160.345
<b>Date:</b>	Fri, 15 Dec 2023	<b>AIC</b>	-2318.689
<b>Time:</b>	20:07:11	<b>BIC</b>	-2314.173
<b>Sample:</b>	0	<b>HQIC</b>	-2316.941
	- 677		
<b>Covariance Type:</b>	opg		

	coef	std err	z	P> z	[0.025	0.975]
<b>sigma2</b>	0.0019	4.69e-05	40.301	0.000	0.002	0.002

<b>Ljung-Box (L1) (Q):</b>	0.03	<b>Jarque-Bera (JB):</b>	1646.47
<b>Prob(Q):</b>	0.85	<b>Prob(JB):</b>	0.00
<b>Heteroskedasticity (H):</b>	0.59	<b>Skew:</b>	0.43
<b>Prob(H) (two-sided):</b>	0.00	<b>Kurtosis:</b>	10.60

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [72]: >>> # Set the significance level (alpha)
alpha = 0.05

p_values = ar_model_zm.pvalues

significant_predictors = p_values[p_values <= alpha]

print("Significant Predictors:")
print(significant_predictors)
```

```
Significant Predictors:
sigma2    0.0
dtype: float64
```



```
In [73]: predicted_original_train_zm = ar_model_zm.predict(start=train_zm_lgp.index[0],
                                                         end=train_zm_lgp.index[-1],
                                                         dynamic=False)
predicted_original_train_zm = np.exp(predicted_original_train_zm)
#predicted_original_ba = np.array(predicted_original_ba)
```

```
In [74]: predicted_original_train_zm = pd.DataFrame(predicted_original_train_zm)

train_val_zm = zm[:int(training_zm_len)]
train_val_zm = pd.merge(train_val_zm, predicted_original_train_zm,
                        left_index=True, right_index=True)
train_val_zm = train_val_zm.rename(columns={'predicted_mean': 'Prediction'})
train_val_zm
```

Out[74]:

	Close	Prediction
Date		
2020-01-02	68.720001	1.000000
2020-01-03	67.279999	68.720001
2020-01-06	70.320000	67.279999
2020-01-07	71.900002	70.320000
2020-01-08	72.550003	71.900002
...	...	...
2022-09-01	81.139999	80.400002
2022-09-02	80.790001	81.139999
2022-09-06	80.019997	80.790001
2022-09-07	81.019997	80.019997
2022-09-08	80.230003	81.019997

677 rows × 2 columns

```
In [75]: train_mse = mean_squared_error(train_val_zm['Close'], train_val_zm['Prediction'])
train_rmse = np.sqrt(train_mse)
print("The MSE of the Trainning set is", train_mse)
print("The RMSE of the Trainning set is", train_rmse)
```

The MSE of the Trainning set is 153.99022624434983  
The RMSE of the Trainning set is 12.409279843905118

```
In [76]: # Create the scaled training data set
test_zm = zm[int(training_zm_len):]
# Log price
test_zm_lgp = np.log(test_zm)

forecast_zm = ar_model_zm.get_forecast(steps=len(test_zm_lgp))

predicted_values_zm = forecast_zm.predicted_mean
predicted_original_zm = np.exp(predicted_values_zm)
predicted_original_zm = np.array(predicted_original_zm)
```

C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:836: ValueWarning: No supported index is available. Prediction results will be given with an integer index beginning at `start`.

```
    return get_prediction_index(
```

```
In [77]: predicted_original_zm = pd.DataFrame(index=test_zm_lgp.index, columns=['Prediction'],
                                         data=predicted_original_zm)

test_zm = pd.merge(test_zm, predicted_original_zm,
                    left_index=True, right_index=True)

test_zm
```

Out[77]:

	Close	Prediction
Date		
2022-09-09	82.620003	80.230003
2022-09-12	84.080002	80.230003
2022-09-13	78.860001	80.230003
2022-09-14	79.589996	80.230003
2022-09-15	80.209999	80.230003
...	...	...
2023-05-05	63.400002	80.230003
2023-05-08	62.930000	80.230003
2023-05-09	63.560001	80.230003
2023-05-10	64.430000	80.230003
2023-05-11	64.449997	80.230003

169 rows × 2 columns

```
In [78]: test_mse = mean_squared_error(test_zm['Close'], test_zm['Prediction'])
test_rmse = np.sqrt(test_mse)
print("The MSE of the Test set is", test_mse)
print("The RMSE of the Test set is", test_rmse)
```

The MSE of the Test set is 87.11304903815056  
The RMSE of the Test set is 9.333437150275913

```
In [79]: # Visualize the data
plt.figure(figsize=(16,6))
plt.title('ARIMA Model (0,1,0) - Zoom (ZM)')
plt.xlabel('Date', fontsize=18)
plt.ylabel('Close Price USD ($)', fontsize=18)
plt.plot(train_val_zm[['Close', 'Prediction']])
plt.plot(test_zm[['Close', 'Prediction']])
plt.legend(['Train_Close', 'Train_Prediction', 'Test_Close', 'Test_Prediction'], loc='upper right')
plt.show()
plt.savefig("images/Zoom_(ZM)_ARIMA_1")
```



In [80]: `pm.arma.auto_arma(train_zm_lgp)`

Out[80]:

	ARIMA
	ARIMA(4,2,1)(0,0,0)[0]

In [81]: `p, d, q = 4, 2, 1`  
`ar_model_zm = ARIMA(train_zm_lgp['Close'], order=(p, d, q))`  
`ar_model_zm = ar_model_zm.fit()`  
`ar_model_zm.summary()`

C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.  
 self.\_init\_dates(dates, freq)  
C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.  
 self.\_init\_dates(dates, freq)  
C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.  
 self.\_init\_dates(dates, freq)

Out[81]: SARIMAX Results

<b>Dep. Variable:</b>	Close	<b>No. Observations:</b>	677
<b>Model:</b>	ARIMA(4, 2, 1)	<b>Log Likelihood</b>	1160.709
<b>Date:</b>	Fri, 15 Dec 2023	<b>AIC</b>	-2309.418
<b>Time:</b>	20:07:20	<b>BIC</b>	-2282.329
<b>Sample:</b>	0	<b>HQIC</b>	-2298.929
	- 677		
<b>Covariance Type:</b>	opg		

	coef	std err	z	P> z	[0.025	0.975]
<b>ar.L1</b>	-0.0289	0.029	-1.000	0.317	-0.085	0.028
<b>ar.L2</b>	-0.0563	0.034	-1.658	0.097	-0.123	0.010
<b>ar.L3</b>	-0.0388	0.038	-1.010	0.312	-0.114	0.037
<b>ar.L4</b>	-0.0607	0.037	-1.654	0.098	-0.133	0.011
<b>ma.L1</b>	-0.9878	0.008	-127.046	0.000	-1.003	-0.973
<b>sigma2</b>	0.0019	4.93e-05	37.914	0.000	0.002	0.002

<b>Ljung-Box (L1) (Q):</b>	0.02	<b>Jarque-Bera (JB):</b>	1609.25
<b>Prob(Q):</b>	0.89	<b>Prob(JB):</b>	0.00
<b>Heteroskedasticity (H):</b>	0.61	<b>Skew:</b>	0.35
<b>Prob(H) (two-sided):</b>	0.00	<b>Kurtosis:</b>	10.53

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [82]: # Set the significance level (alpha)
alpha = 0.05

p_values = ar_model_zm.pvalues

significant_predictors = p_values[p_values <= alpha]

print("Significant Predictors:")
print(significant_predictors)
```

```
Significant Predictors:
ma.L1      0.0
sigma2     0.0
dtype: float64
```

```
In [83]: # p, d, q = 0, 2, 1
# ar_model_zm = ARIMA(train_zm_lgp['Close'], order=(p, d, q))
# ar_model_zm = ar_model_zm.fit()
# ar_model_zm.summary()
```

```
In [84]: # Set the significance level (alpha)
alpha = 0.05

# p_values = ar_model_zm.pvalues

# significant_predictors = p_values[p_values <= alpha]

# print("Significant Predictors:")
# print(significant_predictors)
```

```
In [85]: predicted_original_train_zm = ar_model_zm.predict(start=train_zm_lgp.index[0],
                                                            end=train_zm_lgp.index[-1],
                                                            dynamic=False)
predicted_original_train_zm = np.exp(predicted_original_train_zm)
#predicted_original_ba = np.array(predicted_original_ba)
predicted_original_train_zm = pd.DataFrame(predicted_original_train_zm)

train_val_zm = zm[:int(training_zm_len)]
train_val_zm = pd.merge(train_val_zm, predicted_original_train_zm,
                        left_index=True, right_index=True)
train_val_zm = train_val_zm.rename(columns={'predicted_mean': 'Prediction'})
train_val_zm
```

Out[85]:

	Close	Prediction
Date		
2020-01-02	68.720001	1.000000
2020-01-03	67.279999	569.671822
2020-01-06	70.320000	65.870172
2020-01-07	71.900002	71.198084
2020-01-08	72.550003	72.942304
...	...	...
2022-09-01	81.139999	80.197569
2022-09-02	80.790001	80.716940
2022-09-06	80.019997	80.308033
2022-09-07	81.019997	79.630285
2022-09-08	80.230003	80.550819

677 rows × 2 columns

```
In [86]: train_mse = mean_squared_error(train_val_zm['Close'], train_val_zm['Prediction'])
train_rmse = np.sqrt(train_mse)
print("The MSE of the Trainning set is", train_mse)
print("The RMSE of the Trainning set is", train_rmse)
```

The MSE of the Trainning set is 526.7243910126432

The RMSE of the Trainning set is 22.95047692342456

```
In [87]: # Create the scaled training data set
test_zm = zm[int(training_zm_len):]
# Log price
test_zm_lgp = np.log(test_zm)

forecast_zm = ar_model_zm.get_forecast(steps=len(test_zm_lgp))

predicted_values_zm = forecast_zm.predicted_mean
predicted_original_zm = np.exp(predicted_values_zm)
predicted_original_zm = np.array(predicted_original_zm)
predicted_original_zm = pd.DataFrame(index=test_zm_lgp.index, columns=['Prediction'],
                                     data=predicted_original_zm)

test_zm = pd.merge(test_zm, predicted_original_zm,
                   left_index=True, right_index=True)
test_zm
```

C:\Users\surface\anaconda3\Lib\site-packages\statsmodels\tsa\base\tsa\_model.py:836: ValueWarning: No supported index is available. Prediction results will be given with an integer index beginning at `start`.

```
    return get_prediction_index(
```

Out[87]:

	Close	Prediction
Date		
2022-09-09	82.620003	79.795424
2022-09-12	84.080002	79.410065
2022-09-13	78.860001	78.968458
2022-09-14	79.589996	78.621022
2022-09-15	80.209999	78.253175
...	...	...
2023-05-05	63.400002	36.491449
2023-05-08	62.930000	36.317871
2023-05-09	63.560001	36.145119
2023-05-10	64.430000	35.973188
2023-05-11	64.449997	35.802076

169 rows × 2 columns

```
In [88]: test_mse = mean_squared_error(test_zm['Close'], test_zm['Prediction'])
test_rmse = np.sqrt(test_mse)
print("The MSE of the Test set is", test_mse)
print("The RMSE of the Test set is", test_rmse)
```

The MSE of the Test set is 426.70324320475163

The RMSE of the Test set is 20.656796537816593

```
In [89]: # Visualize the data
plt.figure(figsize=(16,6))
plt.title('ARIMA Model (4,2,1) - Zoom (ZM)')
plt.xlabel('Date', fontsize=18)
plt.ylabel('Close Price USD ($)', fontsize=18)
plt.plot(train_val_zm[['Close', 'Prediction']])
plt.plot(test_zm[['Close', 'Prediction']])
plt.legend(['Train_Close', 'Train_Prediction', 'Test_Close', 'Test_Prediction'], loc='upper right')
plt.show()
plt.savefig("images/Zoom_(ZM)_ARIMA_2")
```



## LSTM

**Step 1: Normalize your data into interval [0,1] (recommend using MinMaxScaler from sklearn.preprocessing)**

```
In [90]: scaler_ba = MinMaxScaler(feature_range=(0,1))
scaled_ba = scaler_ba.fit_transform(ba)
```

```
In [91]: scaler_zm = MinMaxScaler(feature_range=(0,1))
scaled_zm = scaler_zm.fit_transform(zm)
```

**Step 2: Set 80% of the data as training data and 20% of the data as the testing set**

```
In [92]: # Get the number of rows to train the model on
training_ba_len = int(np.ceil( len(scaled_ba) * .8 ))

training_ba_len
```

Out[92]: 677

```
In [93]: # Create the scaled training data set
train_ba = scaled_ba[0:int(training_ba_len), :]
```

```
In [94]: # Get the number of rows to train the model on
training_zm_len = int(np.ceil( len(scaled_zm) * .8 ))

training_zm_len
```

Out[94]: 677

```
In [95]: ▶ # Create the scaled training data set
train_zm = scaled_zm[0:int(training_zm_len), :]
```

### Step 3: Pick a lookback window with 10 timestamps

```
In [96]: ▶ def Split_train_data_with_lookback_window(train_data, timestamps):
    x_train = []
    y_train = []

    for i in range(timestamps, len(train_data)):
        x_train.append(train_data[i-10:i, 0])
        y_train.append(train_data[i, 0])
        if i <= timestamps+1:
            print(x_train)
            print(y_train)
            print()

    # Convert the x_train and y_train to numpy arrays
    x_train, y_train = np.array(x_train), np.array(y_train)

    # Reshape the data
    x_train = np.reshape(x_train, (x_train.shape[0], x_train.shape[1], 1))

    return x_train, y_train
```

```
In [97]: ▶ x_train_ba, y_train_ba = Split_train_data_with_lookback_window(train_ba, 10)

[array([0.94402629, 0.94180795, 0.94568998, 0.95971315, 0.93630163,
        0.95598948, 0.93055776, 0.93174612, 0.94018378, 0.9300823 ])]
[0.9387972920667447]

[array([0.94402629, 0.94180795, 0.94568998, 0.95971315, 0.93630163,
        0.95598948, 0.93055776, 0.93174612, 0.94018378, 0.9300823 ]), array([0.94180795, 0.945689
98, 0.95971315, 0.93630163, 0.95598948,
        0.93055776, 0.93174612, 0.94018378, 0.9300823 , 0.93879729])]
[0.9387972920667447, 0.9077007709514351]
```

```
In [98]: ▶ x_train_zm, y_train_zm = Split_train_data_with_lookback_window(train_zm, 10)

[array([0.01556589, 0.01272856, 0.01871847, 0.02183165, 0.0231124 ,
        0.02325032, 0.02417638, 0.02602853, 0.02431432, 0.0317623 ])]
[0.030126891742712686]

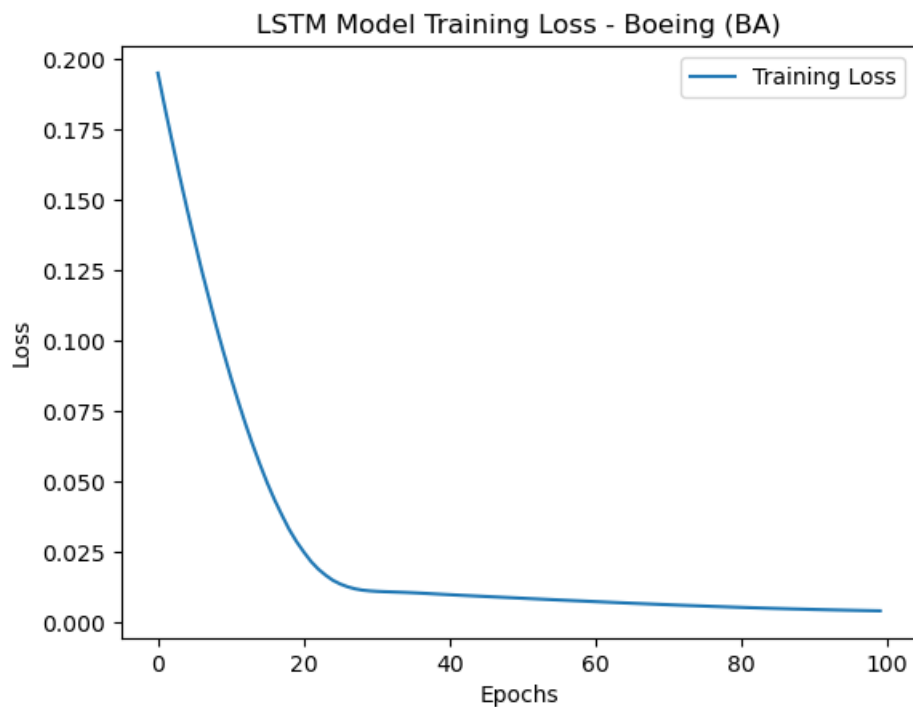
[array([0.01556589, 0.01272856, 0.01871847, 0.02183165, 0.0231124 ,
        0.02325032, 0.02417638, 0.02602853, 0.02431432, 0.0317623 ]), array([0.01272856, 0.018718
47, 0.02183165, 0.0231124 , 0.02325032,
        0.02417638, 0.02602853, 0.02431432, 0.0317623 , 0.03012689])]
[0.030126891742712686, 0.029043184652699228]
```

**Specify the LSTM model with the following parameters:**

```
In [99]: # Build the LSTM model
lstm_model_ba = Sequential()
lstm_model_ba.add(LSTM(4, input_shape= (x_train_ba.shape[1], 1)))
lstm_model_ba.add(Dense(1))
# Compile the model
lstm_model_ba.compile(optimizer='adam', loss='mean_squared_error')
# Train the model
history = lstm_model_ba.fit(x_train_ba, y_train_ba, batch_size=256, epochs=100, verbose=2)
```

```
Epoch 1/100
3/3 - 5s - loss: 0.1949 - 5s/epoch - 2s/step
Epoch 2/100
3/3 - 0s - loss: 0.1821 - 16ms/epoch - 5ms/step
Epoch 3/100
3/3 - 0s - loss: 0.1700 - 31ms/epoch - 10ms/step
Epoch 4/100
3/3 - 0s - loss: 0.1581 - 16ms/epoch - 5ms/step
Epoch 5/100
3/3 - 0s - loss: 0.1467 - 16ms/epoch - 5ms/step
Epoch 6/100
3/3 - 0s - loss: 0.1357 - 16ms/epoch - 5ms/step
Epoch 7/100
3/3 - 0s - loss: 0.1250 - 16ms/epoch - 5ms/step
Epoch 8/100
3/3 - 0s - loss: 0.1150 - 16ms/epoch - 5ms/step
Epoch 9/100
3/3 - 0s - loss: 0.1050 - 31ms/epoch - 10ms/step
Epoch 10/100
3/3 - 0s - loss: 0.0950 - 31ms/epoch - 10ms/step
```

```
In [100]: # Plot the training Loss - Convergence
plt.plot(history.history['loss'], label='Training Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('LSTM Model Training Loss - Boeing (BA)')
plt.legend()
#plt.show()
plt.savefig("images/LSTM Model Training Loss - Boeing (BA)")
```






```
In [101]: # Build the LSTM model
lstm_model_zm = Sequential()
lstm_model_zm.add(LSTM(4, input_shape= (x_train_zm.shape[1], 1)))
lstm_model_zm.add(Dense(1))
# Compile the model
lstm_model_zm.compile(optimizer='adam', loss='mean_squared_error')
# Train the model
history = lstm_model_zm.fit(x_train_zm, y_train_zm, batch_size=256, epochs=100, verbose=2)
```

```
Epoch 1/100
3/3 - 3s - loss: 0.0947 - 3s/epoch - 1s/step
Epoch 2/100
3/3 - 0s - loss: 0.0869 - 16ms/epoch - 5ms/step
Epoch 3/100
3/3 - 0s - loss: 0.0795 - 16ms/epoch - 5ms/step
Epoch 4/100
3/3 - 0s - loss: 0.0725 - 31ms/epoch - 10ms/step
Epoch 5/100
3/3 - 0s - loss: 0.0655 - 16ms/epoch - 5ms/step
Epoch 6/100
3/3 - 0s - loss: 0.0592 - 16ms/epoch - 5ms/step
Epoch 7/100
3/3 - 0s - loss: 0.0532 - 31ms/epoch - 10ms/step
Epoch 8/100
3/3 - 0s - loss: 0.0474 - 24ms/epoch - 8ms/step
Epoch 9/100
3/3 - 0s - loss: 0.0423 - 16ms/epoch - 5ms/step
Epoch 10/100
3/3 - 0s - loss: 0.0373 - 16ms/epoch - 5ms/step
```

```
In [102]: # Plot the training loss
plt.plot(history.history['loss'], label='Training Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('LSTM Model Training Loss - Zoom (ZM)')
plt.legend()
#plt.show()
plt.savefig("images/LSTM Model Trainning Loss - Zoom (ZM)")
```



## Train set predictions and RMSE

```
In [103]: y_train_ba = ba[10:training_ba_len].values
```

```
In [104]: # Get the models predicted price values
train_pred_ba = lstm_model_ba.predict(x_train_ba)
train_pred_ba = scaler_ba.inverse_transform(train_pred_ba)

# Get the root mean squared error (RMSE)
train_rmse_ba = np.sqrt(np.mean(((train_pred_ba - y_train_ba) ** 2)))
train_rmse_ba
```

21/21 [=====] - 1s 2ms/step

Out[104]: 16.034497635163348

```
In [105]: train_val_ba = ba[10:training_ba_len]
train_val_ba['Prediction'] = train_pred_ba
train_val_ba
```

Out[105]:

	Close	Prediction
Date		
2020-01-16	332.000000	303.658203
2020-01-17	324.149994	303.537445
2020-01-21	313.369995	303.064972
2020-01-22	309.000000	301.855530
2020-01-23	317.790009	300.105621
...	...	...
2022-09-01	153.660004	169.883011
2022-09-02	151.820007	168.658813
2022-09-06	152.389999	167.570267
2022-09-07	155.949997	166.774872
2022-09-08	157.789993	166.342865

667 rows × 2 columns

```
In [106]: train_mse = mean_squared_error(train_val_ba['Close'], train_val_ba['Prediction'])
train_rmse = np.sqrt(train_mse)
print("The MSE of the Training set is", train_mse)
print("The RMSE of the Training set is", train_rmse)
```

The MSE of the Training set is 257.10511441205904  
The RMSE of the Training set is 16.034497635163348

```
In [107]: y_train_zm = zm[10:training_zm_len].values
```

```
In [108]: # Get the models predicted price values
train_pred_zm = lstm_model_zm.predict(x_train_zm)
train_pred_zm = scaler_zm.inverse_transform(train_pred_zm)

# Get the root mean squared error (RMSE)
train_rmse_zm = np.sqrt(np.mean(((train_pred_zm - y_train_zm) ** 2)))
train_rmse_zm
```

21/21 [=====] - 1s 2ms/step

Out[108]: 17.990990910365824

```
In [109]: train_val_zm = zm[10:training_zm_len]
train_val_zm['Prediction'] = train_pred_zm
train_val_zm
```

Out[109]:

	Close	Prediction
Date		
2020-01-16	76.110001	80.820297
2020-01-17	75.559998	81.387871
2020-01-21	76.730003	81.700119
2020-01-22	75.540001	82.171432
2020-01-23	74.470001	82.221657
...	...	...
2022-09-01	81.139999	88.512329
2022-09-02	80.790001	87.850410
2022-09-06	80.019997	87.322151
2022-09-07	81.019997	86.778160
2022-09-08	80.230003	86.797424

667 rows × 2 columns

```
In [110]: train_mse = mean_squared_error(train_val_zm['Close'], train_val_zm['Prediction'])
train_rmse = np.sqrt(train_mse)
print("The MSE of the Training set is", train_mse)
print("The RMSE of the Training set is", train_rmse)
```

The MSE of the Training set is 323.67575393686565  
The RMSE of the Training set is 17.990990910365824

## Get Test set and Split to X and y

```
In [111]: ▶ # get test data set and split into X and y
def Split_test_data_with_lookback_window(scaled_data, original_data, training_data_len, timestamps):
    # Create the testing data set
    test_data = scaled_data[training_data_len - 10: , :]
    # Create the data sets x_test and y_test
    x_test = []
    y_test = original_data[training_data_len:].values

    for i in range(timestamps, len(test_data)):
        x_test.append(test_data[i-timestamps:i, 0])

    # Convert the data to a numpy array
    x_test = np.array(x_test)

    # Reshape the data
    x_test = np.reshape(x_test, (x_test.shape[0], x_test.shape[1], 1 ))

    return x_test, y_test
```

```
In [112]: ▶ x_test_ba, y_test_ba = Split_test_data_with_lookback_window(scaled_ba, ba, training_ba_len, 10)
```

```
In [113]: ▶ x_test_zm, y_test_zm, = Split_test_data_with_lookback_window(scaled_zm, zm, training_zm_len, 10)
```

## Prediction

```
In [114]: ▶ # Get the models predicted price values
pred_ba = lstm_model_ba.predict(x_test_ba)
pred_ba = scaler_ba.inverse_transform(pred_ba)

# Get MSE and RMSE
test_mse_ba = np.mean(((pred_ba - y_test_ba) ** 2))
test_rmse_ba = np.sqrt(test_mse_ba)
print("The MSE of the Test set is", test_mse_ba)
print("The RMSE of the Test set is", test_rmse_ba)
```

```
6/6 [=====] - 0s 3ms/step
The MSE of the Test set is 85.99911155398644
The RMSE of the Test set is 9.273570593573245
```

```
In [115]: test_ba = ba[training_ba_len:]
test_ba['Prediction'] = pred_ba
test_ba
```

Out[115]:

	Close	Prediction
Date		
2022-09-09	157.520004	165.988541
2022-09-12	158.720001	165.416107
2022-09-13	147.309998	165.229538
2022-09-14	149.259995	163.906128
2022-09-15	149.779999	163.000687
...	...	...
2023-05-05	198.339996	201.929199
2023-05-08	197.259995	201.254028
2023-05-09	201.880005	200.478256
2023-05-10	200.839996	200.403427
2023-05-11	201.839996	200.212357

169 rows × 2 columns

```
In [116]: # Get the models predicted price values
pred_zm = lstm_model_zm.predict(x_test_zm)
pred_zm = scaler_zm.inverse_transform(pred_zm)

# Get MSE and RMSE
test_mse_zm = np.mean(((pred_zm - y_test_zm) ** 2))
test_rmse_zm = np.sqrt(test_mse_zm)
print("The MSE of the Test set is", test_mse_zm)
print("The RMSE of the Test set is", test_rmse_zm)
```

6/6 [=====] - 0s 0s/step  
The MSE of the Test set is 62.43103856710199  
The RMSE of the Test set is 7.901331442681164

```
In [117]: test_zm = zm[training_zm_len:]
test_zm['Prediction'] = pred_zm
test_zm
```


Out[117]:

	Close	Prediction
Date		
2022-09-09	82.620003	86.598473
2022-09-12	84.080002	86.986015
2022-09-13	78.860001	87.658127
2022-09-14	79.589996	86.942780
2022-09-15	80.209999	86.589653
...	...	...
2023-05-05	63.400002	70.692055
2023-05-08	62.930000	71.057541
2023-05-09	63.560001	71.193703
2023-05-10	64.430000	71.472054
2023-05-11	64.449997	71.893318

169 rows × 2 columns

## Plot the true and the predicted price sequences during the training period plus the testing period

```
In [118]: # Visualize the data
plt.figure(figsize=(16,6))
plt.title('LSTM Model - Boeing (BA)')
plt.xlabel('Date', fontsize=18)
plt.ylabel('Close Price USD ($)', fontsize=18)
plt.plot(train_val_ba[['Close', 'Prediction']])
plt.plot(test_ba[['Close', 'Prediction']])
plt.legend(['Train_Close', 'Train_Prediction', 'Test_Close', 'Test_Prediction'], loc='upper right')
#plt.show()
plt.savefig("images/LSTM Model - Boeing (BA)")
```



```
In [119]: # Visualize the data
plt.figure(figsize=(16,6))
plt.title('LSTM Model - Zoom (ZM)')
plt.xlabel('Date', fontsize=18)
plt.ylabel('Close Price USD ($)', fontsize=18)
plt.plot(train_val_zm[['Close', 'Prediction']])
plt.plot(test_zm[['Close', 'Prediction']])
plt.legend(['Train_Close', 'Train_Prediction', 'Test_Close', 'Test_Prediction'], loc='upper right')
#plt.show()
plt.savefig("images/LSTM Model - Zoom (ZM)")
```

